

TECHNICAL AND DATA SPECIFICATION

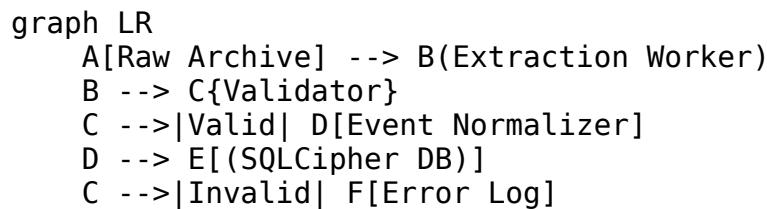
Generated on: 2026-01-30T16:22:34.958Z

FILE: docs/03_ARCHITECTURE/DATA_FLOW_DIAGRAMS.md

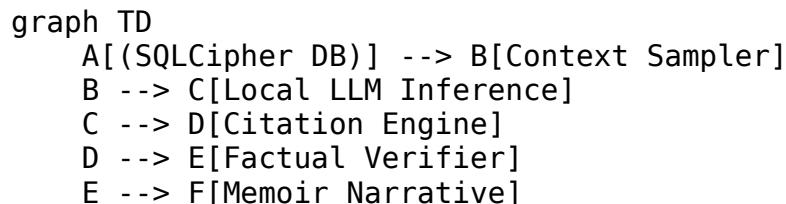
Data Flow Diagrams

High-level visualization of how bits move through Memoir.ai.

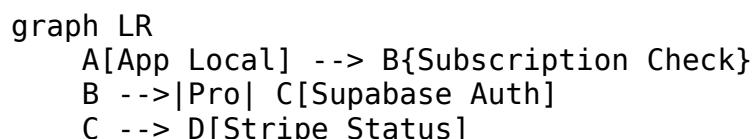
1. Ingestion Flow



2. AI Snapshot Flow



3. Sync Flow (Non-Vault)



FILE: docs/03_ARCHITECTURE/MODULE_MAP.md

Module Map — Memoir.ai

This document maps the logical modules of the Memoir.ai application and their boundaries, as defined in the technical specification.

1. System Components

Frontend (Desktop Client)

- **Env:** Electron (Shell), React (UI Layer).
- **State Management:** Redux (Global Application State).
- **Modules:**
 - **Auth Module:** Handles registration, login, and JWT persistence.
 - **Workspace Module:** Manages the Unified Timeline and Conversation Viewer.
 - **Search Module:** Interfaces with the backend for semantic and keyword queries.
 - **Settings Module:** User profile management and data export controls.

Backend (Local Core)

- **Env:** Node.js, Express.js.
- **Database:** SQLite with SQLCipher (Encrypted Persistence).
- **Modules:**
 - **Ingestion Engine:** Handles multi-source archive parsing and normalization.
 - **Job Runner:** Manages background processing of data imports/exports.
 - **Narrative Engine:** Logic for processing events into summaries and drafts.
 - **Security Layer:** Input validation, sanitization, and encryption management.

2. Component Boundaries

Source	Destination	Protocol	Purpose
Electron Shell	React Frontend	IPC / Context Bridge	System-level integration (File Dialogs, Windows).
React Frontend	Node.js Backend	RESTful APIs (Local)	Data retrieval, Job initiation, Auth requests.
Node.js Backend	SQLite DB	SQL (SQLCipher)	Encrypted data persistence.
Node.js Backend	Local Filesystem	FS API	Raw file ingestion and ZIP export creation.

3. Communication Patterns

- **Request-Response:** Primary pattern for UI interactions (e.g., fetching a timeline range).
- **Asynchronous Jobs:** Background tasks (Parsing, Indexing) reported via progress updates.
- **Event-Driven:** Progress indicators in the frontend responding to job runner state changes.

FILE: docs/03_ARCHITECTURE/SYSTEM_OVERVIEW.md

System Architecture Overview

The structural blueprint of the Memoir.ai "Fortress."

1. Multi-Process Model (Electron)

- **Main Process:** Handles IPC, File System access, and Window lifecycle.
- **Renderer Process:** The "Nebula" UI (React).
- **Worker Processes:** Low-priority background tasks for heavy compute (Ingestion, AI).

2. Storage Layer

- **Vault Partition:** The encrypted SQLCipher database + Media subfolders.
- **Config Partition:** Unencrypted app settings (Theme, Window position).

3. Implementation Stack

See docs/00_META/ADR/ADR_0001_STACK_DECISION.md for full breakdown.

4. Security Philosophy

- **Process Isolation:** The Renderer has no direct access to fs. All data requests go through an IPC bridge with strict validation.
 - **Ephemeral Keys:** Encryption keys are only held in memory for the duration of the session.
-

FILE:

docs/03_ARCHITECTURE/JOB_WORKERS/IMPORT_PIPELINE.md

Import Pipeline Specification — Memoir.ai V1

1. Goal

Provide robust, resumable, and transparent background processing for a single message export format (e.g., CSV/JSON).

2. Pipeline Stages

Step	Component	Action	Output
1. Initiate	Frontend/API	Create <code>import_jobs</code> record, upload	<code>import_jobs</code> record, file in Storage.

Step	Component	Action	Output
2. Process	Worker	Download and parse raw file from Storage.	Structured data stream.
3. Normalize	Worker	Transform raw data into canonical memories schema.	Canonical Memory objects.
4. Ingest	Worker	Batch insert memories into Postgres (linked to library).	New memories records.
5. Finalize	Worker	Update job status to 'complete' or 'failed'.	Finalized import_jobs record.

3. Technical Requirements

- **Resumability:** The worker must handle temporary failures and retry from the last successful batch.
 - **Logging:** Detailed error and status logs must be stored in the `log_output` (JSONB) column.
 - **Progress Tracking:** `progress_percentage` must be updated in real-time for frontend display.
 - **File Isolation:** Raw imports must be stored in user-specific storage paths.
-

FILE:

docs/03_ARCHITECTURE/JOB_RUNNER_SPEC.md

Job Runner Specification

The engine responsible for all heavy background compute in Memoir.ai.

1. Execution Philosophy

- **Local-only:** All jobs run in separate worker threads within the Electron backend.
- **Priority-based:** High-priority jobs (Search Indexing) preempt low-priority jobs (Media Thumbnailing).
- **Non-blocking:** The UI remaining responsive even during 50k+ record imports.

2. Job Lifecycle

1. **Enqueue:** Task added to SQLite-backed queue.
2. **Dispatch:** Worker thread claims the oldest high-priority task.
3. **Active:** Periodic progress events emitted to the UI via IPC.
4. **Completion/Failure:** Results persisted and status updated.

3. Worker Configuration

- **Max Threads:** Configurable (default: CPUs - 1).
- **Memory Throttling:** Jobs pause if application memory usage exceeds 1.5GB to maintain OS stability.
- **Persistence:** Job state is stored in the `import_jobs` table to allow resumption after crash.

4. Specific Runners

- **Parser:** Specialized for `chat.db`, `WhatsApp.zip`, etc.
 - **Indexer:** Updates SQLite FTS5 tables and vector stores.
 - **AI Narrative:** Manages local LLM inference and verification.
-

FILE:

docs/03_ARCHITECTURE/JOB_WORKERS/QUEUE_MODEL.md

Queue Model & Data Structures

Defines how background tasks are represented and managed.

1. Database Schema (`job_queue`)

Field	Type	Description
<code>job_id</code>	UUID	Primary Key
<code>type</code>	String	e.g., <code>PARSE_IMESSAGE</code> , <code>GEN_SNAPSHOT</code>
<code>priority</code>	Integer	0 (Low) to 10 (Critical)
<code>payload</code>	JSON	Job parameters (file paths, UID, etc.)
<code>status</code>	Enum	<code>PENDING</code> , <code>RUNNING</code> , <code>PAUSED</code> , <code>FAILED</code> , <code>DONE</code>
<code>progress</code>	Float	0.0 to 1.0
<code>error_log</code>	Text	Captured stack trace if failed

2. Priority Levels

- **P10 (Urgent)**: Onboarding flow jobs (Initial Vault Setup).
- **P5 (Normal)**: User-triggered imports.
- **P1 (Background)**: Media thumbnailing, long-term integrity scans.

3. Worker Interaction

Workers use a "Pull" model. They query the `job_queue` for the next available PENDING task, mark it as RUNNING within an atomic transaction to prevent double-claiming, and begin execution.

FILE:

[docs/03_ARCHITECTURE/JOBs_WORKERS/RETRIES_IDEMPOTENCY.md](#)

Retries & Idempotency Rules

Ensuring data integrity during failure recovery.

1. Retry Semantic

- **Transient Failures**: (e.g., File system lock) Automatic retry with exponential backoff.
- **Fatal Failures**: (e.g., Disk Full, Permission Denied) Stop and notify user for manual intervention.
- **Max Retries**: Default 3 attempts per job chunk.

2. Idempotency Requirements

- **Parser Idempotency**: Running the same iMessage import twice MUST NOT result in duplicate events. The parser must check the `source_record_id` and `payload_hash` before insertion.
- **Media Ingestion**: Media files are only copied/indexed if their content hash (SHA256) does not already exist in the library.

3. Checkpointing

For large imports (100k+ records), the parser must commit batches to the database every 1,000 records. If the job is interrupted, it resumes from the last successfully committed batch ID.

FILE:

docs/03_ARCHITECTURE/PROVENANCE_VERSIONING/CITATION_SYSTEM.md

Citation System Specification

Ensuring factual traceability between AI narratives and original source evidence.

1. Citation Anatomy

A citation is represented in the database as a link between a **Narrative** and one or more **Events**.

Attribute	Description
citation_id	Unique UUID
target_narrative_id	Foreign Key to Narrative
source_event_ids[]	Array of Event IDs used as supporting evidence
anchor_text	The specific sentence or phrase in the draft being supported
relevance_score	AI-assigned confidence (0.0 - 1.0)

2. UI Representation

In the Snapshot Editor, citations appear as superscript numbers (e.g., `Claim text [1]`).

- **Interaction:** Clicking [1] opens the "Evidence Drawer", displaying the raw message content, timestamp, and source metadata for the linked events.

3. Integrity Rules

- If a source **Event** is deleted, its linked **Citations** are invalidated.
- The UI MUST visually flag narratives with orphaned citations as **UNVERIFIED**.

FILE:

docs/03_ARCHITECTURE/PROVENANCE_VERSIONING/PROVENANCE_MODEL.md

Provenance Model

The chain of custody for every memory record in Memoir.ai.

1. Provenance Meta-Data

Every Event record includes an immutable provenance block:

- **origin_device**: Local hostname where the import occurred.
- **parser_signature**: The version of the parser code used (e.g., imessage-parser-v1.4.2).
- **raw_source_path**: Relative path to the file within the imported archive (e.g., chat.db/message/123).
- **ingestion_timestamp**: When the record was first created in the vault.

2. Verification

During a "Deep Health Check", the system attempts to re-locate the raw record (if the source archive is still present) and verify its **payload_hash** against the current database state.

3. Trust Levels

- **Level 3 (High)**: Directly imported from a platform DB (e.g., chat.db).
 - **Level 2 (Medium)**: Imported from a flattened JSON/CSV export.
 - **Level 1 (Low)**: Manually added or manually edited by the user.
-

FILE:

docs/03_ARCHITECTURE/PROVENANCE_VERSIONING/SNAPSHOT_SYSTEM.md

AI Snapshot System (Citations & Versioning) — Memoir.ai V1

1. AI Snapshot Generation

The system uses an LLM to transform selected memories into a cohesive narrative summary.

Flow

1. **Input Selection**: User selects a set of `memory_ids` or a time range on the timeline.
2. **Prompt Construction**: The backend worker fetches the selected memories and prompts the LLM to generate a summary.
3. **Citation Mapping**: The LLM is instructed to include inline markers (e.g., [1]) that map to specific input memories.
4. **Record Creation**:
 - Creates a new `ai_snapshots` (metadata) record if it's the first in a thread.
 - Creates an `ai_snapshot_versions` record for the narrative content.
 - Creates `snapshot_citations` records linking markers to source `memory_ids`.

2. Data Provenance (Citations)

- **Markers:** Snapshot text contains clickable markers.
- **Source Linking:** Interaction with a marker highlights or navigates the user to the exact source Memory in the timeline.
- **Integrity:** Citations are non-destructive and persistent across versions.

3. Versioning Strategy

- Every "Save" or "Regenerate" operation creates a **new version record**.
 - The UI defaults to the highest `version_number` but allows browsing and reverting to historical versions.
 - This approach ensures user edits do not destroy the original AI output or previous manual drafts.
-

FILE:

docs/03_ARCHITECTURE/PROVENANCE_VERSIONING/SNAPSHOT_VERSIONING_MODEL.md

Snapshot Versioning Model

Managing the evolution of AI-generated narratives.

1. Revision Strategy

Narratives in Memoir.ai are versioned using a simple linear increment (V1, V2, V3).

2. Version Storage

Instead of overwriting, every "Regenerate" operation creates a NEW record in the `narrative_versions` table.

- **Parent ID:** Link to the root Narrative object.
- **Diff:** The system stores the full content for each version (simpler for local AI consumption) rather than incremental diffs.

3. UI Navigation

The "Version Browser" in the Snapshot Editor allows users to:

- **Compare:** Side-by-side view of V[X] and V[Current].
- **Reactivate:** Promote an old version to current.
- **Delete:** Remove bad drafts to save metadata space.

4. Stability Rule

Once a Narrative is "Published" (exported or pinned), the current version is locked. Future changes must create an explicit new version.

FILE:

[docs/03_ARCHITECTURE/SEARCH/INDEXING_STRATEGY.md](#)

Indexing Strategy — Memoir.ai

Indexing is the final stage of the Memoir.ai Ingestion Pipeline, ensuring that raw digital history is transformed into a highly performant, queryable dataset.

1. Indexing Workflow

When a parsing job completes, the system initiates indexing:

1. **Normalization:** Raw data is mapped to the `Events` canonical schema.
2. **Schema-on-Read Optimization:** Common fields are indexed; custom `eventData` JSON fields are made searchable.
3. **Health Check:** Indices are verified for integrity before directing the user to the timeline.

2. Technical Choices

- **Database:** SQLite with SQLCipher.
- **Indices:**
 - B-Tree indices on `id`, `userId`, `dataSourceId`, and `createdAt`.
 - Full-Text Search (FTS) indices (if applicable via SQLite FTS5) for keyword-heavy `eventData`.
- **Batching:** Indexing runs in background workers during the Job Runner phase to prevent main-thread blocking.

3. Security & Performance

- **Encrypted Indices:** All indices are stored within the encrypted SQLCipher vault, maintaining the local-first security posture.
 - **Resumable Jobs:** Indexing jobs are tracked by the Job Runner Dashboard, allowing for recovery if the app is closed during processing.
 - **Latency Target:** Search queries must return results in < 200ms, even for libraries with 50,000+ events.
-

FILE:
docs/03_ARCHITECTURE/SEARCH/QUERY_LANGUAGE_SPEC.md

Query Language Specification

Grammar and reserved keywords for searching the Memoir.ai unified timeline.

1. Syntax Overview

Memoir.ai uses a simplified SQL-like DSL for advanced multi-parameter queries, which are translated into SQLite FTS5 queries.

2. Keywords & Facets

Facet	Description	Example
from:	Filter by sender/author	from: "Mom"
to:	Filter by recipient	to: "Dad"
source:	Filter by platform	source: imessage
date:	Specific date or range	date: 2023-01-01..2023-12-31
type:	Event category	type: media or type: text
has:	Content attributes	has: link or has: attachment

3. Logical Operators

- **AND:** Implicit (e.g., from: Mom source: whatsapp)
- **OR:** Explicit (e.g., from: Mom OR from: Dad)
- **NOT:** Prefix with - (e.g., source: whatsapp -has: link)

4. Semantic Search

Prefixing a query with ~ triggers a semantic/vector search instead of a keyword match.

- Example: ~ "When was the last time we talked about the trip?"

5. Implementation Note

The frontend input component parses these tokens and highlights them for the user. The backend then constructs the WHERE clause dynamically, ensuring all inputs are sanitized to prevent SQL injection.

FILE:

docs/03_ARCHITECTURE/SEARCH/SEARCH_OVERVIEW.md

Search Subsystem Overview

1. Hybrid Architecture

Combines Lexical (SQLite FTS5) and Semantic (Vector) search.

2. Relevancy Scoring

Final score = $(0.4 * \text{LexicalScore}) + (0.6 * \text{SemanticScore})$.

3. Privacy

Vector embeddings are computed locally and stored in the encrypted vault.

FILE:

docs/03_ARCHITECTURE/PERFORMANCE/CACHING_STRATEGY.md

Caching Strategy

Optimizing the "Private Time Machine" for local performance.

1. UI State Caching

- **Redux Persist:** Stores UI preferences but NOT user data (to avoid unencrypted spill).
- **Component Memoization:** Heavy use of React .memo for the Timeline feed entries.

2. Database Caching

- **SQLite Page Cache:** Configured to 128MB to keep hot indices in memory.
- **Prepared Statements:** Re-used for recurring search filters to avoid parse overhead.

3. Media Caching

- **Thumbnail generation:** During import, the app generates low-res WebP thumbnails stored in the `vault/.thumb/` directory.
 - **Action:** The UI displays thumbnails while loading full-res assets from the encrypted blob store.
-

FILE:
docs/03_ARCHITECTURE/PERFORMANCE/LARGE_IMPORTS_ST
RATEGY.md

Large Imports Strategy

Handling 10GB+ archives without application degradation.

1. The "Chunking" Pattern

The ingestion engine never reads a full file into memory. It uses Node.js streams to process records in chunks of 500.

2. Resource Throttling

- **CPU:** Workers are restricted to N - 1 cores to ensure the Main process remains responsive for UI input.
- **Priority:** Ingestion jobs are marked with IDLE_PRIORITY so they yield to real-time AI generation or search requests.

3. Interruption Recovery

- **Checkpointing:** After every 1,000 records, the progress is committed to SQLite.
 - **Action:** If the app crashes or the computer shuts down, the import resumes from the last successful checkpoint.
-

FILE:
docs/03_ARCHITECTURE/PERFORMANCE/PERF_BUDGETS.md

Performance Budgets

SLA-style targets for local interaction.

Metric	Budget (P95)	Context
Vault Unlock	< 1.0s	PBKDF2 Hashing -> DB Open
Timeline Scroll	60 FPS	Jitter-free vertical movement
Global Search	< 250ms	Keyword search across 50k events
Semantic Search	< 1.5s	Vector embedding -> Lookup
Snapshot Draft	< 8.0s	Model inference for 1

Metric	Budget (P95)	Context
Cold Start	< 2.0s	chapter App launch -> Splash screen

FILE: docs/04_DATA/CANONICAL_DATA_MODEL.md

Canonical Data Model

The unified schema for all personal history.

1. The Event Entity

The core atom of Memoir.ai.

Field	Type	Description
<code>event_id</code>	UUID	Primary Key
<code>source_id</code>	UUID	FK to the data source
<code>timestamp</code>	ISO8601	Precision to milliseconds
<code>platform</code>	Enum	iMessage, WhatsApp, Email, etc.
<code>content_raw</code>	Text	Raw unformatted body
<code>metadata</code>	JSONB	Platform-specific headers

2. The Person Entity

Extracted participants.

Field	Type	Description
<code>person_id</code>	UUID	Primary Key
<code>display_name</code>	String	User-provided or extracted
<code>identities</code>	JSONB	Map of phone numbers/emails

3. Relationships

- **Event -> Person:** Many-to-Many via `participants` table.
- **Event -> Snapshot:** Linked via `citations`.

FILE: docs/04_DATA/ENTITY_DICTIONARY.md

Entity Dictionary

Detailed semantic definitions for project terminology.

Core Entities

- **DataSource:** A connection to a remote/local archive (e.g., "Troy's iPhone iMessage").
- **Attachment:** A binary file (Image, Video, Audio) linked to an Event.
- **Thread:** A chronological cluster of messages between a set of participants.
- **NarrativeBlock:** A single paragraph or section of an AI Snapshot.

Field Semantics

- `normalized_at`: The UTC time when the record was finalized in the vault.
 - `entropy_score`: Internal metric for the information density of a message (used by AI sampler).
 - `citation_anchor`: The exact range of characters in a narrative linked to a source.
-

FILE:
docs/04_DATA/DATA_VALIDATION/DEDUPE_MERGE_RULES.md

Deduplication & Merge Rules

Managing overlapping records from multiple data sources.

1. Deduplication (Exact Match)

Records are merged if they meet ALL of the following criteria:

- **Timestamp:** Identical within ± 1 second.
- **Participants:** Exact match of resolved entity IDs.
- **Payload Hash:** SHA-256 of the normalized text content is identical.

2. Fuzzy Merging (Overlapping Events)

If records originate from different sources (e.g., an iMessage and a WhatsApp export) but describe the same event:

- **Rule:** If the timestamp is within ± 30 seconds and the text similarity is $> 90\%$, group them as a "Consolidated Event".
- **Metadata:** Retain BOTH `source_ids` and `raw_payloads` specifically for provenance verification.

3. Media Deduplication

- **Rule:** All media is indexed by SHA-256 content hash.
- **Action:** If a photo is sent in iMessage and later imported from a local folder, it is stored only ONCE. The database creates two `event_media` links pointing to the same media entry.

4. Conflict Resolution

If two sources provide conflicting metadata (e.g., different "Read" timestamps):

- **Priority:** Higher-integrity databases (e.g., live `chat.db`) override flattened ZIP exports.
 - **Conservative Approach:** If priority is equal, keep BOTH values under a `metadata_conflicts` field.
-

FILE:

[docs/04_DATA/DATA_VALIDATION/IMPORT_VALIDATION_RULES.md](#)

Import Validation Rules

Strict quality gates for data entering the Memoir.ai vault.

1. Schema Validation

Every record must conform to the Event schema before it is inserted into the encrypted datastore.

- **Required Fields:** `event_id`, `timestamp_utc`, `source_id`, `event_type`.
- **Constraint:** `timestamp_utc` must be a valid ISO 8601 string or Unix Epoch.

2. Source Integrity Checks

- **MBOX:** Must have valid headers (`From:`, `Date:`) and a non-empty body.
- **SQLite (iMessage/WhatsApp):** The table schema must match the expected version for that platform's parser.
- **JSON Exports:** Must conform to the specific platform's current export structure (e.g., Instagram v3 Export).

3. Data Cleansing

- Remove NULL characters or malformed byte sequences that could crash indexers.
- Trim whitespace from participant names.
- Resolve relative media paths to absolute paths within the temporary import sandbox.

4. Rejection Policy

If > 10% of a batch fails structural validation, the entire ImportJob is paused, and the user is prompted for a "Deep Scan" or "Repair" of the source file.

FILE:

docs/04_DATA/DATA_VALIDATION/NORMALIZATION_RULES.md

Normalization Rules

Transforming diverse data formats into the "Memoir.ai Canonical Schema".

1. Participant Normalization

- **Mapping:** Map platform-specific handles (e.g., +123456789, user_123, mom@email.com) to a single `person_id`.
- **Primary Name:** Use the richest metadata available (e.g., Apple Contacts name overrides raw phone number).

2. Timestamp Standardization

- All internal storage is in **UTC**.
- The `UI_DISPLAY_TIME` is calculated on the fly using the `vault_timezone` stored in the user profile.

3. Event Type Mapping

Raw Source Type	Canonical <code>event_type</code>
iMessage SMS / WhatsApp Chat	MESSAGE
Instagram Post / FB Wall	SOCIAL_POST
Gmail / Outlook	EMAIL
JPG / PNG (No message)	PHOTO_EVENT

4. Attachment Handling

- Attachments are extracted and assigned a unique `media_id`.
 - If an event is a "Group Chat", the normalization layer creates links to all participants in the `event_participants` junction table.
-

Export Schema (JSON)

Ensuring data portability without proprietary lock-in.

1. Export Scope

The export includes all normalized events, narratives, and participant data in a machine-readable format.

2. Event Structure

```
{  
  "export_metadata": {  
    "vault_id": "uuid",  
    "exported_at": "timestamp",  
    "version": "1.0.0"  
  },  
  "events": [  
    {  
      "id": "event_uuid",  
      "timestamp": "2023-01-01T12:00:00Z",  
      "type": "MESSAGE",  
      "source": "imessage",  
      "author": "Alice",  
      "content": "Hello world!",  
      "attachments": ["media_id_1.jpg"],  
      "provenance": "chat.db/message/123"  
    }  
  ]  
}
```

3. Narrative Structure

AI drafts include citations linked back to event IDs within the same export bundle.

```
{  
  "narratives": [  
    {  
      "id": "narrative_uuid",  
      "title": "Summer Trip",  
      "body": "We went to the beach [1].",  
      "citations": [  
        { "id": 1, "evidence_ids": ["event_uuid_456"] }  
      ]  
    }  
  ]  
}
```

FILE:
docs/04_DATA/EXPORT_FORMATS/EXPORT_ZIP_LAYOUT.md

Export ZIP Layout

Organization of the human-readable data package.

1. Directory Structure

```
Memoir_Export_[Date]/
└── data/
    ├── events.json
    ├── narratives.json
    └── library_metadata.json
└── media/
    ├── images/
    ├── videos/
    └── originals/
└── docs/
    ├── Narratives_Markdown/ (User-readable mirror)
    ├── README_EXPORT_MAP.md
    └── LICENSE_DATA_SOVEREIGNTY.txt
```

2. Media Organization

- Media is stored in the `media/` folder, indexed by their SHA256 hash (e.g., `media/images/a1b2c3d4...jpg`).
- The `events.json` file uses these relative paths for easy reconstruction by external tools.

3. Narrative Mirroring

To ensure users can read their memoirs without special software, every `Narrative` is also exported as a standalone `.md` file in `docs/Narratives_Markdown/`.

4. Integrity

The root level contains a `checksums.txt` file listing the MD5/SHA256 hashes of every file in the export bundle.

FILE:
docs/04_DATA/EXPORT_FORMATS/JSON_SCHEMAS/provenance.schema.json

```
{  
  "$schema": "http://json-schema.org/draft-07/schema#",
```

```
"title": "Provenance Export Schema",
"type": "object",
"properties": {
  "export_id": { "type": "string", "format": "uuid" },
  "citations": {
    "type": "array",
    "items": {
      "type": "object",
      "properties": {
        "source_event_id": { "type": "string" },
        "target_fragment": { "type": "string" },
        "confidence_score": { "type": "number" }
      }
    }
  }
}
```

FILE:

docs/04_DATA/EXPORT_FORMATS/JSON_SCHEMAS/snapshot.schema.json

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "title": "Snapshot Export Schema",
  "type": "object",
  "properties": {
    "snapshot_id": {
      "type": "string"
    },
    "version": {
      "type": "integer"
    },
    "narrative_body_markdown": {
      "type": "string"
    },
    "generated_at": {
      "type": "string",
      "format": "date-time"
    }
  }
}
```

FILE:

docs/04_DATA/EXPORT_FORMATS/JSON_SCHEMAS/timeline.schema.json

```
{  
    "$schema": "http://json-schema.org/draft-07/schema#",  
    "title": "Timeline Export Schema",  
    "type": "object",  
    "properties": {  
        "timeline_id": {  
            "type": "string"  
        },  
        "entries": {  
            "type": "array",  
            "items": {  
                "$ref": "#/definitions/entry"  
            }  
        }  
    },  
    "definitions": {  
        "entry": {  
            "type": "object",  
            "properties": {  
                "event_id": {  
                    "type": "string"  
                },  
                "timestamp": {  
                    "type": "string"  
                },  
                "platform": {  
                    "type": "string"  
                },  
                "content": {  
                    "type": "string"  
                }  
            }  
        }  
    }  
}
```

FILE:
docs/04_DATA/SUPABASE/SCHEMAS/00_SCHEMA_OVERVIEW.md

Schema Overview

High-level map of the Supabase backend tables.

1. The Metadata Hub

While the "Private Bits" live in the Vault, the Supabase backend tracks:

- Subscription visibility.
- Sync health.
- Cross-device metadata (e.g., "Troy's MacBook" vs "Troy's PC").

2. Entity Map

See the individual .sql files in this directory for the full relational schema.

- **Auth Layer:** 01_AUTH_WORKSPACES.sql
 - **Job Layer:** 02_SOURCES_IMPORT_JOBS.sql
 - **Snapshot Layer:** 06_SNAPSHOTS VERSIONS.sql
-

FILE:
docs/04_DATA/SUPABASE/SCHEMAS/01_AUTH_WORKSPACES.sql

```
-- 01_AUTH_WORKSPACES.sql
CREATE TABLE auth_profiles (
    id UUID REFERENCES auth.users PRIMARY KEY,
    full_name TEXT,
    avatar_url TEXT
);

CREATE TABLE workspaces (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    owner_id UUID REFERENCES auth.users,
    slug TEXT UNIQUE,
    settings JSONB DEFAULT '{}'
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/02_SOURCES_IMPORT_JOBS.sql

```
-- 02_SOURCES_IMPORT_JOBS.sql
CREATE TABLE data_sources (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    source_type TEXT NOT NULL,
    config JSONB
);

CREATE TABLE import_jobs (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    source_id UUID REFERENCES data_sources,
    status TEXT,
    error_log TEXT
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/03_PARTICIPANTS_CONVERSATIONS_MESSAGES.sql

```
-- 03_PARTICIPANTS_CONVERSATIONS_MESSAGES.sql
CREATE TABLE conversation_threads (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    platform TEXT,
    external_id TEXT,
    metadata JSONB
);

CREATE TABLE participants (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    display_name TEXT,
    contact_info JSONB
);

CREATE TABLE message_records (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    thread_id UUID REFERENCES conversation_threads,
    sender_id UUID REFERENCES participants,
    content_text TEXT,
    sent_at TIMESTAMPTZ
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/04_MEDIA_ATTACHMENT
S.sql

```
-- 04_MEDIA_ATTACHMENTS.sql
CREATE TABLE media_assets (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    file_path TEXT,
    mime_type TEXT,
    size_bytes BIGINT,
    metadata JSONB
);

CREATE TABLE message_attachments (
    message_id UUID REFERENCES message_records,
    media_id UUID REFERENCES media_assets,
    PRIMARY KEY (message_id, media_id)
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/05_TIMELINE_ENTRIES.sql

```
-- 05_TIMELINE_ENTRIES.sql
CREATE TABLE timeline_events (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    event_type TEXT,
    occurred_at TIMESTAMPTZ,
    reference_id UUID, -- Link to message, photo, etc.
    summary_text TEXT
);

CREATE INDEX idx_timeline_workspace_time ON timeline_events
(workspace_id, occurred_at DESC);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/06_SNAPSHOTS_VERSION
S.sql

```
-- 06_SNAPSHOTS VERSIONS.sql
CREATE TABLE snapshot_records (
```

```
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
workspace_id UUID REFERENCES workspaces,
version_id INTEGER DEFAULT 1,
title TEXT,
narrative_blob_ref TEXT -- Path in Supabase Storage or UUID
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/07_CITATIONS_PROVENANCE.sql

```
-- 07 CITATIONS PROVENANCE.sql
CREATE TABLE citations (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    snapshot_id UUID REFERENCES snapshot_records,
    source_event_id UUID,
    confidence_score FLOAT,
    context_snippet TEXT
);
```

FILE:

docs/04_DATA/SUPABASE/SCHEMAS/08_TAGS_NOTES_BOOKMARKS.sql

```
-- 08 TAGS NOTES BOOKMARKS.sql
CREATE TABLE user_tags (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    label TEXT,
    color TEXT
);

CREATE TABLE tagged_items (
    tag_id UUID REFERENCES user_tags,
    item_id UUID, -- Link to Event, Participant, etc.
    item_type TEXT,
    PRIMARY KEY (tag_id, item_id)
);
```

FILE:
docs/04_DATA/SUPABASE/SCHEMAS/09_EXPORT_DELETE_AUDIT.sql

```
-- 09_EXPORT_DELETE_AUDIT.sql
CREATE TABLE audit_logs (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    user_id UUID REFERENCES auth.users,
    action TEXT,
    metadata JSONB,
    created_at TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE export_history (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces,
    status TEXT,
    download_url TEXT,
    expired_at TIMESTAMPTZ
);
```

FILE:
docs/04_DATA/SUPABASE/RLS_POLICIES/00_RLS_OVERVIEW.md

RLS Overview

Row Level Security (RLS) is the primary defense for multi-tenant isolation in Memoir.ai's cloud metadata layer.

1. Core Principle

- **No Global Reads:** All tables must have RLS enabled.
- **Ownership Check:** Policies MUST verify that `auth.uid() == owner_id` or equivalent.

2. Policy Matrix

Table	SELECT	INSERT/UPDATE
workspaces	owner	owner
import_jobs	workspace owner	workspace owner
entitlements	user	service_role only

3. Enforcement

The Supabase anon key is restricted via these policies. The `service_role` key is used only by internal webhooks (e.g., Stripe) to bypass RLS for administrative updates.

FILE:

```
docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_AUDIT_LOGS.sql
-- RLS for audit logs
ALTER TABLE audit_logs ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Users can only view their own audit logs"
ON audit_logs FOR SELECT
USING (auth.uid() = user_id);

CREATE POLICY "System can insert audit logs"
ON audit_logs FOR INSERT
WITH CHECK (true); -- Usually restricted to service_role via API
```

FILE: docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_MEDIA.sql

```
-- RLS for media tables
ALTER TABLE media_assets ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Users can only access media from their workspaces"
ON media_assets FOR SELECT
USING (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));

CREATE POLICY "Users can upload media to their workspaces"
ON media_assets FOR INSERT
WITH CHECK (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));
```

FILE:

```
docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_MESSAGES.sql
-- RLS for messages (Cloud metadata layer only)
ALTER TABLE cloud_message_metadata ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Users can only read their own message metadata"
ON cloud_message_metadata FOR SELECT
```

```
USING (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));
```

FILE:

docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_SNAPSHOTS.sql

```
-- RLS for snapshots
ALTER TABLE snapshot_records ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Owners manage their snapshots"
ON snapshot_records FOR ALL
USING (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));
```

FILE:

docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_SOURCES_IMPO
RTS.sql

```
-- RLS for sources and imports
ALTER TABLE data_sources ENABLE ROW LEVEL SECURITY;
ALTER TABLE import_jobs ENABLE ROW LEVEL SECURITY;

CREATE POLICY "Users manage their own data sources"
ON data_sources FOR ALL
USING (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));

CREATE POLICY "Users manage their own import jobs"
ON import_jobs FOR ALL
USING (auth.uid() IN (
    SELECT owner_id FROM workspaces WHERE id = workspace_id
));
```

FILE:

docs/04_DATA/SUPABASE/RLS_POLICIES/RLS_WORKSPACES.s
ql

```
-- RLS for workspaces
ALTER TABLE workspaces ENABLE ROW LEVEL SECURITY;
```

```
CREATE POLICY "Owners can see their workspaces"
ON workspaces FOR SELECT
USING (auth.uid() = owner_id);

CREATE POLICY "Owners can create workspaces"
ON workspaces FOR INSERT
WITH CHECK (auth.uid() = owner_id);
```

FILE: docs/04_DATA/SUPABASE/MIGRATIONS/0001_init.sql

```
-- 0001_init.sql
-- Initial migration for Memoir.ai backend.

CREATE TABLE workspaces (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    owner_id UUID REFERENCES auth.users(id),
    name TEXT NOT NULL,
    created_at TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE import_jobs (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces(id),
    status TEXT CHECK (status IN ('PENDING', 'RUNNING', 'COMPLETED',
'FAILED')),
    source_type TEXT NOT NULL,
    metadata JSONB DEFAULT '{}',
    created_at TIMESTAMPTZ DEFAULT now()
);

CREATE TABLE entitlements (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    user_id UUID REFERENCES auth.users(id),
    tier TEXT DEFAULT 'FREE',
    stripe_customer_id TEXT,
    updated_at TIMESTAMPTZ DEFAULT now()
);
```

FILE:

docs/04_DATA/SUPABASE/MIGRATIONS/0002_add_snapshot_versioning.sql

```
-- 0002_add_snapshot_versioning.sql
```

```
ALTER TABLE import_jobs ADD COLUMN version_id UUID;
```

```
CREATE TABLE snapshot_metadata (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces(id),
    version_key TEXT NOT NULL,
    checksum TEXT,
    created_at TIMESTAMPTZ DEFAULT now()
);
```

FILE:

docs/04_DATA/SUPABASE/MIGRATIONS/0003_add_citations.sql
-- 0003_add_citations.sql

```
CREATE TABLE provenance_sync (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    workspace_id UUID REFERENCES workspaces(id),
    event_id_map JSONB, -- Mapping local Event UUIDs to sync refs
    last_synced_at TIMESTAMPTZ DEFAULT now()
);
```

FILE: docs/04_DATA/SUPABASE/SEEDS/seed_dev_minimal.sql

```
-- Seed Dev Minimal
-- Insert test user (Fake UUID)
INSERT INTO auth.users (id, email) VALUES ('f47ac10b-58cc-4372-a567-0e02b2c3d479', 'test@memoir.ai');

-- Insert initial workspace
INSERT INTO workspaces (owner_id, name) VALUES ('f47ac10b-58cc-4372-a567-0e02b2c3d479', 'Default Vault');

-- Insert Free Tier entitlement
INSERT INTO entitlements (user_id, tier) VALUES ('f47ac10b-58cc-4372-a567-0e02b2c3d479', 'FREE');
```

FILE: docs/05_APIS/API_OVERVIEW.md

API Overview

The internal communication layer between the React UI and the Node.js/Electron backend.

1. Protocol

- **IPC (Inter-Process Communication):** Main communication channel for sensitive vault operations.
- **RESTful Endpoints (Local):** Used for non-sensitive data retrieval (e.g., UI preferences, non-encrypted metadata).

2. Global Constants

- **Base Local URL:** `http://localhost:port/api/v1`
- **Authentication:** All requests must include the `X-Vault-Token` generated during the unlock flow.

3. Core Controllers

- `/api/vault`: Initialize, Unlock, Lock, Change Passphrase.
- `/api/timeline`: Fetch events, Filter, Search.
- `/api/jobs`: Enqueue import, Monitor progress, Cancel jobs.
- `/api/narratives`: CRUD operations for Snapshots and Drafts.
- `/api/billing`: Subscription status, Stripe session triggers.

4. Response Format

```
{  
  "success": true,  
  "data": { ... },  
  "error": null,  
  "meta": {  
    "timestamp": "ISO-8601",  
    "requestId": "uuid"  
  }  
}
```

FILE: docs/05_APIS/AUTHORIZATION_MODEL.md

Authorization Model

Determining access rights within the local-first application.

1. The Master Key

Full access to the vault is strictly gated by the **Passphrase**.

- **Status: Locked:** Only the `/auth` routes are available. All Event and Narrative requests return `403 Forbidden`.
- **Status: Unlocked:** The backend holds the SQLCipher master key in memory (sanitized on lock). All local API calls are permitted.

2. Role-Based Access Control (RBAC)

While it is an individual-user app, internal RBAC applies to premium feature access:

- **Free:** Restricted imports (max 2 sources), no AI narratives.
- **Pro:** Unlimited sources, full AI Snapshot generation.
- **Team/Shared:** Future capability for multiple IDs within a single multi-participant vault.

3. Entitlement Checks

Every request to the Jobs or AI layers performs a check against the `subscription_status` cached in the secure local preferences.

FILE: docs/05_APIS/ERROR_TAXONOMY.md

Error Taxonomy

Standardized error codes and user-facing recovery messages.

1. Standard Error Object

```
{  
  "code": "VAULT_LOCKED",  
  "message": "The vault must be unlocked to perform this action.",  
  "recovery": "Navigate to the lock screen and enter your  
passphrase.",  
  "severity": "FATAL"  
}
```

2. Error Categories

Code Prefix	Category	Description
AUTH_	Authentication	Passphrase mismatch, session expiry.
VALT_	Vault Integrity	Corruption, missing DB file, storage full.
IMPT_	Import Engine	Corrupt archive, unknown format, parser crash.
AI_	AI Services	Token limit, hallucination guard fail, model offline.
BILL_	Commercials	Payment failed, feature gated.

3. Logging Policy

- **UI:** Display the message and recovery action in a "Nebula Red" toast or modal.
 - **Backend logs:** Record the code and detailed stack trace locally in `logs/error.log` for troubleshooting.
-

FILE: docs/05_APIS/ENDPOINTS/API_ENDPOINTS.md

API Endpoints Specification

Public and internal REST endpoints for the Memoir.ai sync service.

1. Auth & Workspaces

- `POST /v1/auth/login`: Exchange Supabase credentials for a session.
- `GET /v1/workspaces`: List all associated vaults for the user.
- `POST /v1/workspaces`: Initialize a new cloud record for a local vault.

2. Billing & Entitlements

- `GET /v1/billing/status`: Fetch current subscription tier and usage counts.
- `POST /v1/billing/checkout`: Generate a Stripe Checkout URL.
- `GET /v1/billing/invoices`: List recent invoice metadata.

3. Sync & Jobs

- `POST /v1/sync/heartbeat`: Update the cloud with local health metrics.
 - `GET /v1/jobs/:id/status`: Check the progress of an async cloud verification job.
-

FILE: docs/05_APIS/ENDPOINTS/BILLING_API.md

Billing API Specification

Endpoints

- `GET /v1/billing/plans`: Fetch available subscription tiers.
- `POST /v1/billing/subscribe`: Initiate Stripe checkout.
- `DELETE /v1/billing/cancel`: Stop recurring payments.

Responses

Standard JSON response with [ERROR_TAXONOMY.md](#) codes.

FILE: docs/05_APIS/ENDPOINTS/EXPORT_DELETE_API.md

Export/Delete API Specification

1. Data Export

- `POST /v1/export/request`: Trigger a full vault export.
- `GET /v1/export/status/:id`: Check status of a ZIP generation.
- `GET /v1/export/download/:id`: Get a temporary signed URL for the export.

2. Account Deletion

- `POST /v1/account/delete`: Initiate full account and data purge.
 - `POST /v1/account/delete/confirm`: Securely confirm deletion with MFA/Password.
-

FILE: docs/05_APIS/ENDPOINTS/IMPORTS_API.md

Imports API Specification

1. Import Jobs

- `POST /v1/imports/jobs`: Start a new ingestion process.
- `GET /v1/imports/jobs`: List recent import history.
- `GET /v1/imports/jobs/:id`: Detailed status and error logs.

2. Sources

- `GET /v1/imports/sources`: List configured data connections.
 - `DELETE /v1/imports/sources/:id`: Remove a source connection.
-

FILE: docs/05_APIS/ENDPOINTS/MEDIA_API.md

Media API Specification

1. Asset Management

- `GET /v1/media`: Search media assets with filters.
 - `GET /v1/media/:id/url`: Get a temporary URL for high-res viewing.
 - `GET /v1/media/:id/thumbnail`: Get low-res preview.
-

FILE: docs/05_APIS/ENDPOINTS/SEARCH_API.md

Search API Specification

Endpoints

- `POST /v1/search/query`: Submit a hybrid search request.
- `GET /v1/search/history`: Retrieve recent query strings.

Relevancy Parameters

- `alpha`: Balance between Lexical and Semantic.
 - `limit`: Max records to return.
-

FILE: docs/05_APIS/ENDPOINTS/SNAPSHOTS_API.md

Snapshots API Specification

1. Generation

- `POST /v1/snapshots/generate`: Request AI narrative generation.
- `GET /v1/snapshots`: List available snapshots.
- `GET /v1/snapshots/:id`: Retrieve full snapshot body and citations.

2. Iteration

- `PATCH /v1/snapshots/:id`: Save manual edits to a narrative.
 - `POST /v1/snapshots/:id/regenerate`: Trigger a new version with updated prompts.
-

FILE: docs/05_APIS/ENDPOINTS/TIMELINE_API.md

Timeline API Specification

1. Discovery

- `GET /v1/timeline`: Fetch a page of timeline events.
 - `GET /v1/timeline/count`: Total event count for progress bars.
 - `GET /v1/timeline/stats`: Distribution of events over time/platform.
-

FILE: docs/05_APIS/EVENT_SCHEMAS/EVENT_SCHEMAS.md

Event Schemas

The shape of messages sent over the IPC and Sync channels.

1. VAULT_LOCKED Event

Triggered when the user locks their database or the timeout expires.

```
{  
  "type": "VAULT_LOCKED",  
  "payload": {  
    "vault_id": "UUID",  
    "timestamp": "ISO8601",  
    "reason": "USER_ACTION | TIMEOUT"  
  }  
}
```

2. IMPORT_PROGRESS Event

Emitted by the Ingestion Worker.

```
{  
  "type": "IMPORT_PROGRESS",  
  "payload": {  
    "job_id": "UUID",  
    "percent_complete": 45,  
    "current_file": "chat.db"  
  }  
}
```

FILE: docs/05_APIS/EVENT_SCHEMAS/export_ready.event.json

```
{  
  "$schema": "http://json-schema.org/draft-07/schema#",  
  "title": "Export Ready Event Schema",  
  "type": "object",  
  "properties": {  
    "export_id": {  
      "type": "string"  
    },  
    "download_url": {  
      "type": "string"  
    },  
    "expiry": {  
      "type": "string",  
      "format": "date-time"  
    }  
  }  
}
```

```
        }
    }
}
```

FILE: docs/05_APIS/EVENT_SCHEMAS/import_job.event.json

```
{
    "$schema": "http://json-schema.org/draft-07/schema#",
    "title": "Import Job Event Schema",
    "type": "object",
    "properties": {
        "job_id": {
            "type": "string"
        },
        "status": {
            "enum": [
                "PENDING",
                "RUNNING",
                "COMPLETED",
                "FAILED"
            ]
        },
        "records_processed": {
            "type": "integer"
        }
    }
}
```

FILE:

docs/05_APIS/EVENT_SCHEMAS/snapshot_generated.event.json

```
{
    "$schema": "http://json-schema.org/draft-07/schema#",
    "title": "Snapshot Generated Event Schema",
    "type": "object",
    "properties": {
        "snapshot_id": {
            "type": "string"
        },
        "version": {
            "type": "integer"
        },
        "uri": {
            "type": "string"
        }
    }
}
```

FILE: docs/05_APIS/WEBHOOKS/STRIPE_WEBHOOKS.md

Stripe Webhooks

1. Lifecycle Events

- `checkout.session.completed`: Upgrades user to Pro.
- `customer.subscription.deleted`: Downgrades user to Free.
- `invoice.payment_failed`: Triggers notification flow for payment issues.

2. Security

All webhooks MUST verify the `stripe-signature` header using the signing secret.

FILE: docs/05_APIS/WEBHOOKS/SUPABASE_TRIGGER.md

Supabase Triggers

1. Database Events

- `on_user_created`: SQL trigger to insert default workspace and free entitlement.
- `on_workspace_deleted`: Cleanup trigger to purge linked metadata.

2. Real-time Listeners

The application listens for `INSERT` on `import_jobs` to refresh the UI progress state.

FILE: docs/05_APIS/WEBHOOKS/WEBHOOKS_SPEC.md

Webhooks Specification

External callbacks handled by the Memoir.ai backend.

1. Stripe Checkout Complete

- **Route:** POST `/webhooks/stripe/checkout-complete`
- **Action:** Update `entitlements` table to reflect new subscription tier.

2. Stripe Invoice Paid

- **Route:** POST `/webhooks/stripe/invoice-paid`

- **Action:** Reset usage meters for the new billing cycle.

3. Supabase Auth Signup

- **Route:** POST /webhooks/auth/on-signup (Database Trigger)
 - **Action:** Create default "Solo" workspace record.
-